

ECE 485/585
Fall 2021
Final Project Description

Your team is responsible for the simulation of a memory controller capable serving the shared last level cache of a four core 3.2 GHz processor employing a single memory channel. The system uses a relaxed consistency (XC) model. The memory channel is populated by a single-ranked 8GB PC4-25600 DIMM (constructed with memory chips organized as x8 devices with a 2KB page size and 24-24-24 timing). There is no ECC.

The memory controller must exploit bank parallelism and an open page policy. Assume the DIMM is already initialized and all banks are in the pre-charged state. Do not model additive latency or CRC. The controller should maintain a queue of 16 outstanding memory requests.

You can use any programming language (e.g. C, C++, Java) or hardware description language (e.g. SystemVerilog) you like to implement the simulation. Note that you are not doing a synthesizable design, though you must keep track of DRAM cycle counts.

Timing constraints

Your memory controller simulation must correctly schedule DRAM commands to comply with the following DDR-4 timing values. Unless otherwise indicated, all values are given in DRAM clock cycles.

Parameter	Value
tRC	76
tRAS	52
tRRD_L	6
tRRD_S	4
tRP	24
tRFC	350ns
CWL (tCWD)	20
tCAS (CL)	24
tRCD	24
tWR	20
tRTP	12
tCCD_L	8
tCCD_S	4
tBURST	4
tWTR_L	12
tWTR_S	4
REFI	7.8μs

The meaning and use of these parameters can be found in the course lecture slides and in datasheets for DDR4 devices such as this [Micron DDR4 device](#). Do not assume that this device is in any way similar in speed, size, organization, or timing to the device(s) in this assignment. The datasheet is provided only so that you can see how to interpret the parameters given above.

Input format

Your program must accept input trace files with the following format:

<time> <operation> <hexadecimal address>

where <time> is the time of the request in (absolute) *CPU* clock cycles from the beginning of the trace execution, <operation> is either 0, 1, or 2 to indicate the type of memory request being made (data read, data write, instruction fetch, respectively), and <hexadecimal address> is the hexadecimal address being referenced by the operation. Fields are separate by one or more spaces or tabs. All addresses will be 8-byte aligned. The last level caches employ critical word first and early restart.

For example:

```
30 2 0x01FF97000
31 2 0x01FF97080
32 0 0x10FFFFFF00
40 1 0x10FFFFFF80
```

Output format

Your program must generate as output a text file with a trace of all DRAM commands issues using the following format:

<time> <DRAM command>

where <time> is the time in (absolute) elapsed *CPU* cycles when the DRAM command is issued and <DRAM command> is one of the following:

```
ACT <bank group> <bank> <row>
PRE <bank group> <bank>
RD <bank group> <bank> <column>
WR <bank group> <bank> <column>
REF
```

and <bank group>, <bank>, <row>, and <column> are hexadecimal values. Note that you may not use RD or WR with auto pre-charge. All fields are separated by one or more spaces or tabs. For readability, you should display the time in a long fixed-width field

so that the commands align. You may want to display the commands in fixed-width fields as well.

For example:

```
100  PRE 0 0
200  ACT 0 0 03FF
300  RD  0 0 EF
```

GitHub

Your team must use GitHub to host your code. If you are unfamiliar with GitHub there are several very good tutorials at guides.github.com.

Grading

The project is worth 100 points. Your grade will be based upon:

- External specification
- Completeness of the solution (adherence to the requirements above)
- Correctness of the solution
- Quality and readability of the project report (e.g. specifications, design decisions)
- Validity of design decisions
- Quality of implementation
- Structure and clarity of design
- Readability
- Maintainability
- Testing
- Presentation of results and ability to explain your project and the underlying concepts and issues.

Note that your score can be lower (possibly significantly) if you have bugs: incorrect results, DRAM timing violations, etc. or your code is poorly written (difficult to understand, unreliable, difficult to maintain, hard to re-use).

There are a variety of policies a memory controller can employ in an attempt to improve performance (usually bandwidth, though sometimes worst-case latency or power).

You can earn extra credit by simulating more aggressive scheduling policies. If you do this, you must provide a run-time switch to enable/disable them. Examples of optional scheduling policies include: adaptive open page policy and out-of-order scheduling (e.g. to prioritize accesses to open row in the same bank or reads before writes).

You can also receive extra credit for implementing refresh operations, further exploration of the design space (e.g. implementing and measuring performance of different out-of-order scheduling policies, request buffer size).

Your simulation does not need to model or support mode register programming, low power modes, or additive latency.

Scheduling policies

The in-order scheduling, open page policy processes requests in order. It doesn't explicitly close a page unless the queue is non-empty and the next in-order reference is to a different row in the same bank.

First ready, first access scheduling still processes requests in the order received. However, it can schedule individual DRAM commands from more than one in-flight request so that, for example, after issuing an activate command for one memory reference it can issue the activate command to another bank/bank group for another request before issuing the read command for the earlier reference.

Out-of-order policies can process memory requests out of order to optimize the DRAM command schedule. There are many possibilities. The simplest would be to allow a memory reference to a bank that could accommodate a request immediately even if it arrived after another memory reference (to another bank) that can't be started. Another might be to schedule a memory read ahead of one that arrived earlier if it's to the currently open row in the same bank. Another might schedule a read ahead of a write (provided they are not to the same address!). Out-of-order policies must incorporate an "aging"/timeout process so that a request isn't indefinitely passed over by later requests.

Academic honesty

You are free to consult the literature (e.g. papers in conferences and journals) for ideas provided you cite all sources in your final report/presentation. You cannot make use of any existing code or have anyone outside your team write code for you. Doing so will result in a code of conduct violation complaint and a score of 0 on the assignment for all team members.

Competition

All teams doing more than the minimum in-order, no access scheduling, open page policy can participate in a competition to achieve the best bandwidth on a benchmark that I will provide.

Teams with top three performance will be recognized and earn extra credit. The team with the top-performing controller will win a small prize and be able to smugly mock their unworthy competitors.

Presentation and demonstration

Each team will demo their final project in a 30 minute time slot during finals week. You will provide a brief written report of your project as well as all documentation, testcases, and source code along with either a makefile (with default target) or README

file with instructions for building and running your project. These should be packaged in a single zip file and uploaded to D2L 24 hours prior to your demo slot.

Verification (testing) is important. Be sure to include a description of your test strategy and your testcases in your report.

Checkpoints

Because of the complexity of the project it's important that you get started early and make steady progress throughout the quarter. I've devised a series of early and intermediate checkpoints to help you with this. It has the added advantage that any major misunderstandings can be caught early when you can still correct them. Your team's performance during these checkpoints is part of your grade. The dates for each checkpoint will be shown on D2L. Your team will need demonstrate to the course TA or instructor the capability indicated. I may revise the assignment and modify details of the checkpoints as the term progresses.

1. -- Working code to read and parse a input trace file (the name of which is specified by the user)
 - Debugging code (that be enabled/disabled at either runtime or compile time) to display the parsed values for each line
 - Use of github repository
2. -- Diagram showing how address bits are mapped to required fields (e.g. row, column, bank, bank group).
3. -- Have working code to insert items into the queue(s), stub to remove items from queue(s), and maintain correct (simulated) time. What information needs to be maintained in the queue?
4. -- Demonstrate code generates correct output format for DRAM commands.